

Q1_IFT3165

February 17, 2019

```
In [ ]: # Importing dataset
import urllib.request
import _pickle as pickle
import gzip
import os
import numpy as np
%matplotlib inline
import matplotlib.pyplot as plt

path = 'http://deeplearning.net/data/mnist'
mnist_filename_all = 'mnist.pkl'
local_filename = os.path.join("./", mnist_filename_all)
urllib.request.urlretrieve(
    "{}/{}.gz".format(path, mnist_filename_all), local_filename+'.gz')
tr,va,te = pickle.load(gzip.open(local_filename+'.gz','r'), encoding='latin1')
np.save(open(local_filename+'.npy','wb'), (tr,va,te))

In [ ]: # Split dataset

validRatio = 0.1
mnist_npy = np.load("mnist.pkl.npy")
train_images = mnist_npy[0,0]
train_labels = mnist_npy[0,1]
valid_images = mnist_npy[1,0]
valid_labels = mnist_npy[1,1]
test_images = mnist_npy[2,0]
test_labels = mnist_npy[2,1]

def one_hot_vectorize(vector):
    one_hot_matrix = np.zeros((vector.shape[0], np.max(vector) + 1))
    one_hot_matrix[np.arange(vector.shape[0]), vector] = 1
    return one_hot_matrix

train_labels = one_hot_vectorize(train_labels)
valid_labels = one_hot_vectorize(valid_labels)
test_labels = one_hot_vectorize(test_labels)

In [ ]: # Building basic modules
import numpy as np
```

```

from abc import ABC, abstractmethod
from enum import Enum

class Function(ABC):
    @abstractmethod
    def fprop(self, input):
        pass

    @abstractmethod
    def bprop(self, grads_output):
        pass

class InitMethod(Enum):
    ZERO = 1
    NORMAL = 2
    GLOROT = 3
    XAVIER = 4

class LinearLayer(Function):
    def __init__(self, in_n_nodes, out_n_nodes, init_method=InitMethod.GLOROT):
        self.in_n_nodes = in_n_nodes
        self.out_n_nodes = out_n_nodes
        self.init_method = init_method

    def reset(self):
        if self.init_method == InitMethod.ZERO:
            self.weights = np.zeros((self.in_n_nodes, self.out_n_nodes))
        elif self.init_method == InitMethod.NORMAL:
            self.weights = np.random.normal(size=(self.in_n_nodes, self.out_n_nodes))
        elif self.init_method == InitMethod.GLOROT:
            bound = np.sqrt(6.0/(1.0*
                                (self.in_n_nodes
                                 + self.out_n_nodes)))
            self.weights = np.random.uniform(-bound,
                                              bound,
                                              size=(self.in_n_nodes,
                                                    self.out_n_nodes))
        elif self.init_method == InitMethod.XAVIER:
            bound = 1/np.sqrt(self.in_n_nodes)
            self.weights = np.random.uniform(-bound,
                                              bound,
                                              size=(self.in_n_nodes,
                                                    self.out_n_nodes))

        self.bias = np.zeros(self.out_n_nodes)
        self.iterations = 0
        self.zero_grads()

    def zero_grads(self):

```

```

        self.grad_weights = np.zeros_like(self.weights)
        self.grad_bias = np.zeros_like(self.bias)
        self.grad_input = np.zeros((1, self.in_n_nodes))

    def fprop(self, input):
        self.input = input
        output = self.input.dot(self.weights) + self.bias
        assert(output.shape == (self.input.shape[0], self.out_n_nodes))
        return output

    def bprop(self, grad_output):
        self.grad_weights = (self.input.T.dot(grad_output)) / self.input.shape[0]
        assert(self.grad_weights.shape == self.weights.shape)

        self.grad_bias = np.mean(grad_output, axis=0)
        assert(self.grad_bias.shape == self.bias.shape)

        self.grad_input = grad_output.dot(self.weights.T)
        assert(self.grad_input.shape == self.input.shape)

        return self.grad_input

    def update(self, learning_rate):
        self.weights -= learning_rate * self.grad_weights
        self.bias -= learning_rate * self.grad_bias

class ReLU(Function):
    def fprop(self, input):
        self.input = input
        output = np.maximum(self.input, 0)
        assert(output.shape == self.input.shape)
        return output

    def bprop(self, grad_output):
        grad_input = (self.input > 0) * grad_output
        return grad_input

class Softmax():
    @staticmethod
    def fprop(input):
        exp = np.exp(input - input.max(axis=1, keepdims=True))
        output = exp / exp.sum(axis=1, keepdims=True)
        return output

class CrossEntropyLoss(Function):
    def fprop(self, input, target):
        self.epsilon = 1e-10
        self.target = target

```

```

self.input = Softmax.fprop(input)
assert(np.allclose(self.input.sum(axis=1, keepdims=True),
                    np.ones_like(self.target.shape[0])))
logs = -np.log(self.input[np.arange(input.shape[0]),
                              self.target.argmax(axis=1)]+self.epsilon)
self.losses = logs.sum()
return (self.losses) / input.shape[0]

def bprop(self):
    self.grad_input = copy.deepcopy(self.input)
    self.grad_input[np.arange(self.input.shape[0]), self.target.argmax(axis=1)] -= 1
    return self.grad_input

```

In []: *# Building NN*

```

import copy

class NN(object):
    def __init__(self, hidden_dims=(2,2), init_method=InitMethod.GLOROT):
        dims = (784, *hidden_dims)
        self.fcs = []
        self.layers = []
        for i_dim in range(len(dims) - 1):
            fc = LinearLayer(dims[i_dim], dims[i_dim + 1], init_method)
            relu = ReLU()
            self.fcs.append(fc)
            self.layers.append(fc)
            self.layers.append(relu)

        last_linear = LinearLayer(dims[-1], 10, init_method)
        self.fcs.append(last_linear)
        self.layers.append(last_linear)

        self.lossFunc = CrossEntropyLoss()

        self.initialize_weights()

    def initialize_weights(self):
        for fc in self.fcs:
            fc.reset()

    def zero_grads(self):
        for fc in self.fcs:
            fc.zero_grads()

    def forward(self, input):
        x = input
        for layer in self.layers:
            x = layer.fprop(x)

```

```

        return x

def loss(self, prediction, labels):
    return self.lossFunc.fprop(prediction, labels)

def softmax(self, input):
    output = Softmax.fprop(input)
    return output

def backward(self):
    grad = self.lossFunc.bprop()
    for layer in reversed(self.layers):
        grad = layer.bprop(grad)

def update(self, learning_rate):
    for fc in self.fcs:
        fc.update(learning_rate)

In [ ]: def finite_gradient_check(mlp, x, y, k=None, i=None):
    if not k:
        k = np.random.randint(1, 5)
    if not i:
        i = np.random.randint(0, 5)
    eps = 1/float(k*10**i)

    weights = copy.deepcopy(mlp.fcs[1].weights)
    n_weights = min(10, weights.shape[1])
    estimated_grads = np.zeros_like(weights[:n_weights])

    for i_idx in range(n_weights):
        for j_idx in range(weights.shape[0]):
            mlp.fcs[1].weights[i_idx, j_idx] += eps

            output = mlp.forward(x)
            loss_peps = mlp.loss(output, y)

            mlp.fcs[1].weights[i_idx, j_idx] -= 2*eps

            output = mlp.forward(x)
            loss_neps = mlp.loss(output, y)

            estimated_grads[i_idx, j_idx] = (loss_peps - loss_neps)/(2.0*eps)

    mlp.fcs[1].weights = weights
    output = mlp.forward(x)
    loss = mlp.loss(output, y)
    mlp.backward()
    # print("Gradient estimé")

```

```

# print(estimated_grads)

# print("Gradient par backpropagation")
# print(mlp.fcs[1].grad_weights)
return estimated_grads, mlp.fcs[1].grad_weights

In [ ]: import matplotlib.axes as ax
np.random.seed(6) # Adjust seed if no straight line is shown in the graph
mlp = NN((2,2),init_method=InitMethod.GLOROT)
maxDiff = []
xaxis = []

for i in range(3,6):
    for k in range(1,6):
        estimatedGrads, bpropGrads = finite_gradient_check(mlp,
                                                            train_images[0].reshape(1,-1),
                                                            train_labels[0].reshape(1,-1),
                                                            k, i)
        maxDiff.append(np.amax(np.absolute(estimatedGrads-bpropGrads)))
        xaxis.append(1/float(k*10**i))

xaxis.reverse()
maxDiff.reverse()

plt.plot(xaxis,maxDiff)

plt.xlabel('$\epsilon$ value')
plt.ylabel('Maxium gradient difference')

In [ ]: def batchLoader(dataset, batchSize):
    images, labels = dataset
    batches = []
    eof = False
    i = 0
    while eof == False:
        if (i+batchSize)>= len(images):
            batches.append((images[i:],labels[i:]))
            eof = True
        else:
            batches.append((images[i:i+batchSize],labels[i:i+batchSize]))
            i += batchSize
    return batches

In [ ]: def trainEpoch(mlp, dataset, batchSize, learningRate):
    batches = batchLoader(dataset,batchSize)
    loss_avg = 0
    acc_avg = 0
    for i, (images, labels) in enumerate(batches):

```

```

        output = mlp.forward(images.reshape(len(images), -1))
        loss = mlp.loss(output, labels)
        loss_avg += loss
        mlp.backward()
        mlp.update(learningRate)
        mlp.zero_grads()
        correct = (np.argmax(output, axis=1) == np.argmax(labels, axis=1)).sum()
        acc_avg += correct
        if (i+1) % 100 == 0:
            print('Step [{}/{}], Loss: {:.4f}({:.4f}), Accuracy: {:.3f}({:.3f})'
                  .format(i+1, len(batches), loss, loss_avg / (i + 1),
                          correct * 100 / batchSize,
                          acc_avg * 100 / ((i + 1) * batchSize)))
    return loss_avg / (i + 1), acc_avg * 100 / ((i + 1) * batchSize)

In [ ]: def modelEval(mlp, dataset):
        images, labels = dataset
        output = mlp.forward(images.reshape(len(images), -1))
        loss = np.mean(mlp.loss(output, labels))
        rightPred = 0
        total = 0
        for i, example in enumerate(output):
            if np.argmax(example) == np.argmax(labels[i]):
                rightPred += 1
            total += 1
        print('Validation Loss: {:.4f}, Accuracy: {:.4f} '
              .format(loss, rightPred * 100 / total))
        return loss, rightPred / total

In [ ]: # ZERO INIT
        mlp = NN((512, 512), init_method=InitMethod.ZERO)
        losses_zero_init = []
        for epoch in range(10):
            print("Epoch [{}/10]".format(epoch))
            loss = trainEpoch(mlp, (train_images, train_labels), 20, 0.0005)
            losses_zero_init.append(loss)

In [ ]: # NORMAL INIT
        mlp = NN((512, 512), init_method=InitMethod.NORMAL)
        losses_normal_init = []
        for epoch in range(10):
            print("Epoch [{}/10]".format(epoch))
            loss = trainEpoch(mlp, (train_images, train_labels), 20, 0.0005)
            losses_normal_init.append(loss)

In [ ]: # GLOROT INIT
        mlp = NN((512, 512), init_method=InitMethod.GLOROT)
        losses_glorot_init = []
        for epoch in range(10):

```

```

    print("Epoch [{}/10]".format(epoch))
    loss = trainEpoch(mlp, (train_images, train_labels), 20, 0.0005)
    losses_glorot_init.append(loss)

In [ ]: # XAVIER INIT
mlp = NN((512,512), init_method=InitMethod.XAVIER)
losses_xavier_init = []
for epoch in range(10):
    print("Epoch [{}/10]".format(epoch))
    loss = trainEpoch(mlp, (train_images, train_labels), 20, 0.0005)
    losses_xavier_init.append(loss)

In [ ]: #Plotting losses
x = np.arange(1, len(losses_glorot_init)+1)
plt.subplot(1,2,1)
plt.plot(x,[losses_zero_init[i][0] for i in range(len(losses_zero_init))])
plt.plot(x,[losses_normal_init[i][0] for i in range(len(losses_normal_init))])
plt.plot(x,[losses_glorot_init[i][0] for i in range(len(losses_glorot_init))])

plt.title("Loss curve as a function of the epoch number")
plt.xlabel("Epoch number")
plt.ylabel("Loss")

plt.subplot(1,2,2)
plt.plot(x,[losses_zero_init[i][1] for i in range(len(losses_zero_init))])
plt.plot(x,[losses_normal_init[i][1] for i in range(len(losses_normal_init))])
plt.plot(x,[losses_glorot_init[i][1] for i in range(len(losses_glorot_init))])

plt.title("Accuracy curve as a function of the epoch number")
plt.xlabel("Epoch number")
plt.ylabel("Accuracy")
plt.legend(["Zero","Normal","Glorot"], title = "Initialization method")

fig = plt.gcf()
fig.set_size_inches(12, 6)

In [ ]: # test main
model = NN((512,512), init_method=InitMethod.GLOROT)
trainDataset = [train_images,train_labels]
validDataset = [valid_images,valid_labels]
testDataset = [test_images,test_labels]
nbEpoch = 30
batchSize = 128
learningRate = 0.05

train_loss = []
train_acc = []
valid_loss = []

```



```

valid_accuracy = []

for epoch in range(1,nbEpoch+1):
    print("Epoch [{}/{}]".format(epoch, nbEpoch))
    # train
    loss, acc = trainEpoch(model, trainDataset, batchSize, learningRate)
    train_loss.append(loss)
    train_acc.append(acc)

    # valid
    loss, acc = modelEval(model, validDataset)
    valid_loss.append(loss)
    if epoch % 17 == 0: #(epoch>1 and np.abs(acc-valid_accuracy[-1])<0.01):
        learningRate /=10
        print('New learning rate : {}'.format(learningRate))
    valid_accuracy.append(acc)

# test
print('Test performances: \n')
loss, acc = modelEval(model, testDataset)

x = np.arange(len(train_loss))
plt.plot(x,train_loss)
plt.plot(x,valid_loss)

In [ ]: x = np.arange(len(train_loss))
plt.subplot(1,2,1)
plt.plot(x,train_loss)
plt.plot(x,valid_loss)
plt.xlabel("Epoch number")
plt.ylabel("Loss")

plt.subplot(1,2,2)
plt.plot(x,train_acc)
plt.plot(x,100*np.asarray(valid_accuracy))
plt.xlabel("Epoch number")
plt.ylabel("Accuracy")

plt.legend(["Training", "Validation"])

fig = plt.gcf()
fig.set_size_inches(12, 6)

```